Machine Learning vs. Deep Learning: Real-World Applications

Duy Nguyen

Introduction

Machine learning (ML) and deep learning (DL) both use data to discover patterns and make predictions, but they differ in how they represent information and the resources they typically require. Traditional ML can be especially effective for structured business records, while deep learning is often useful for complex, high-dimensional inputs such as images. Comparing customer churn prediction and medical image analysis illustrates why selecting the right method matters.

Traditional Machine Learning: Customer Churn Prediction

A real-world application is a telecommunications provider predicting which customers may cancel service. A company can train a support vector machine (SVM) on historical records containing contract length, monthly charges, tenure, service usage, support interactions, and whether each customer previously left. Once trained and validated, the model can flag accounts for a retention team to review. This is a practical example of supervised learning because past customer outcomes provide labels.

An SVM is suitable because these inputs are largely structured and can be represented as a manageable set of numerical or encoded categorical features. With appropriate scaling, feature selection, and validation, an SVM can find a useful boundary between customers likely and unlikely to churn. It may perform well even when the dataset is not large enough to justify training a complex neural network. Its predictions should still be evaluated for precision, recall, and fairness, since falsely flagging customers can waste retention resources.

Deep learning is not necessarily unsuitable, but it is less compelling as a first choice here. A multilayer neural network may need more tuning and computation without delivering a meaningful improvement over a well-designed traditional model on tabular data. If the company later incorporates large volumes of call transcripts or other unstructured information, deep learning could become more attractive.

Deep Learning: Medical Image Recognition

A second real-world application is the use of convolutional neural networks (CNNs) to analyze medical images, such as chest X-rays, and identify patterns that may warrant a radiologist's attention. CNN-based computer-aided systems can learn from labeled images and produce a probability or visual indication of potentially abnormal findings. Such systems support clinical review rather than replacing professional diagnosis.

Deep learning is suitable because an image contains thousands or millions of pixel values whose spatial relationships matter. CNN layers learn increasingly complex visual features, from edges and textures to combinations of shapes. This reduces the need for experts to manually define every visual feature. With representative training data, careful external validation, and clinical oversight, a CNN can help prioritize studies or provide a second assessment.

A traditional ML approach is generally less suitable for this raw-image task because it typically relies on manually engineered descriptors such as texture measurements or predefined shapes. These features may miss subtle variations and can be brittle across imaging devices or patient populations. Traditional ML remains viable when strong handcrafted image features are available or data are limited; however, for large and diverse image collections, CNNs often offer a more flexible way to learn useful representations. Deep learning also introduces risks, including dataset bias, limited explainability, and potentially costly training, so performance must be tested carefully before clinical use.

Comparison and Conclusion

The key distinction is not that deep learning is always better. Customer churn prediction begins with a compact set of meaningful business variables, making traditional ML an efficient and defensible starting point. Medical image recognition requires interpreting complex spatial patterns, making deep learning more appropriate. In both cases, the best model is the one that meets the actual problem requirements and is evaluated on reliable, representative data. From a data analytics perspective, I would begin with the simplest credible baseline, compare results using task-appropriate metrics, and consider the costs, transparency, and consequences of errors before deployment.

References

Course lesson. (n.d.). Machine learning vs. deep learning [Assignment learning material].

Esteva, A., Kuprel, B., Novoa, R. A., Ko, J., Swetter, S. M., Blau, H. M., & Thrun, S. (2017). Dermatologist-level classification of skin cancer with deep neural networks. Nature, 542, 115–118. https://doi.org/10.1038/nature21056

Huang, B., Kechadi, M. T., & Buckley, B. (2012). Customer churn prediction in telecommunications. Expert Systems with Applications, 39(1), 1414–1425. https://doi.org/10.1016/j.eswa.2011.08.024